

## Group 28:

Student ID: 24280041 24280046

### This document contains:

#### 1. Script and Screen Shot of Outputs for:

- Step 1: Ingestion Script
- Step 2: Raw Tables in Hive
- Step 3: Star Schema
- Step 4: Transformation
- Step 5: Queries

#### 2. Hive DDL from Docker Terminal With all scripts working and outputs in the same order as above

**Step 1:** Used these commands to ingest data manually first before making it into .sh file and enabling it.

Used these commands to make directory in HDFS

```
/bin/sh: 1: docker: not found
# hdfs dfs -mkdir -p /raw/logs
hdfs dfs -mkdir -p /raw/metadata
hdfs dfs -mkdir -p /user/hive/warehouse
# # # hdfs dfs -ls /raw
Found 2 items
drwxr-xr-x  - root supergroup          0 2025-03-11 07:47 /raw/logs
drwxr-xr-x  - root supergroup          0 2025-03-11 07:47 /raw/metadata
# ls -l /home/
total 20
-rwxr-xr-x 1 root  root    431 Mar  4 08:44 content_metadata.csv
drwxr-x--- 2 ubuntu ubuntu 4096 Aug  1 2024 ubuntu
```


Then moved files manually to the new made directories in HDFS

```
-rwxr-xr-x 1 root root 431 Mar 4 08:44 content_metadata.csv
drwxr-x--- 2 ubuntu ubuntu 4096 Aug 1 2024 ubuntu
-rwxr-xr-x 1 root root 9764 Mar 4 08:44 user_logs.csv
# hdfs dfs -put /home/user_logs.csv /raw/logs/
# hdfs dfs -put /home/content_metadata.csv /raw/metadata/
# hdfs dfs -ls /raw/logs
Found 1 items
-rw-r--r-- 1 root supergroup 9764 2025-03-11 07:52 /raw/logs/user_logs.csv
# hdfs dfs -ls /raw/metadata
Found 1 items
-rw-r--r-- 1 root supergroup 431 2025-03-11 07:52 /raw/metadata/content_metadata.csv
# hive
SLF4J: Class path contains multiple SLF4J bindings.
```

RAM 6.55 GB CPU 0.25% Disk: 6.66 GB used (limit 1006.85 GB)

Terminal New version ava

## Created tables in hive to test

 Docker Debug brings the tools you need to debug your container with one click.  
Requires a paid Docker subscription. [Learn more.](#)

Upgrade

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS user_logs (
>   user_id STRING,
>   event_type STRING,
>   event_timestamp STRING,
>   event_details STRING
> )
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS TEXTFILE
> LOCATION '/raw/logs/';
OK
Time taken: 3.843 seconds
hive> CREATE EXTERNAL TABLE IF NOT EXISTS content_metadata (content_id STRING,title STRING,genre STRING,release_year INT)
```

## Step 2:

### External Table for user\_logs (Partitioned)

Since the logs table should be partitioned by (year, month, day), run:

```
CREATE EXTERNAL TABLE IF NOT EXISTS raw_user_logs (
```

```
    user_id INT,
```

```
    content_id INT,
```

```
    action STRING,
```

```
    timestamp STRING,
```

```
    device STRING,
```

```
    region STRING,
```

```
    session_id STRING
```

```
)  
PARTITIONED BY (year INT, month INT, day INT)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE  
LOCATION '/raw/logs';
```

### **Add Partitions**

Since Hive doesn't automatically detect partitions for external tables, we need to add them manually:

```
ALTER TABLE raw_user_logs ADD PARTITION (year=2025, month=03, day=11) LOCATION  
'/raw/logs/2025/03/11/';
```

### **External Table for content\_metadata**

```
CREATE EXTERNAL TABLE IF NOT EXISTS raw_content_metadata (  
    content_id INT,  
    title STRING,  
    category STRING,  
    length INT,  
    artist STRING  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE  
LOCATION '/raw/metadata';
```

## **Step 3:**

**SQL For Tables and Star Schema. (Hive DDL given at the end)**

**Star schema** as follows:

- **Fact Table:** fact\_user\_activity (user interactions with content)
- **Dimension Tables:**
  - dim\_users (user-related details)
  - dim\_content (content metadata)
  - dim\_sessions (session-related details)

### **dim\_users Table**

```
CREATE TABLE dim_users STORED AS PARQUET AS
SELECT DISTINCT user_id FROM raw_user_logs;
```

### **dim\_content Table (Ignoring NULL columns)**

sql

CopyEdit

```
CREATE TABLE dim_content STORED AS PARQUET AS
SELECT DISTINCT content_id, title, category, artist
FROM raw_content_metadata
WHERE content_id IS NOT NULL;
```

### **dim\_sessions Table**

sql

CopyEdit

```
CREATE TABLE dim_sessions STORED AS PARQUET AS
SELECT DISTINCT session_id, device, region
FROM raw_user_logs;
```

## **Step 4:**

### **Fact Table & Load Data Using INSERT OVERWRITE**

sql

CopyEdit

```
CREATE TABLE fact_user_activity (  
    user_id INT,  
    content_id INT,  
    session_id STRING,  
    action STRING,  
    event_timestamp STRING,  
    year INT,  
    month INT,  
    day INT  
) STORED AS PARQUET;  
  
INSERT OVERWRITE TABLE fact_user_activity  
SELECT user_id, content_id, session_id, action, `timestamp`, year, month, day  
FROM raw_user_logs;
```

## Step 5:

### Query 1:

Counts distinct users per region, per month.

```
SELECT  
    f.year,  
    f.month,  
    s.region,  
    COUNT(DISTINCT f.user_id) AS monthly_active_users  
FROM fact_user_activity f  
JOIN dim_sessions s ON f.session_id = s.session_id  
WHERE f.year = 2025 -- Example filter  
GROUP BY f.year, f.month, s.region
```

ORDER BY f.year, f.month, monthly\_active\_users DESC;

```
MapReduce Total cumulative CPU time: 2 seconds 190 msec
Ended Job = job_1741695021252_0006
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 4.56 sec HDI
Stage-Stage-3: Map: 1 Reduce: 1 Cumulative CPU: 2.19 sec HDI
Total MapReduce CPU Time Spent: 6 seconds 750 msec
OK
2025      3      US      80
2025      3      APAC    79
2025      3      EU      72
2025      3      region  0
Time taken: 48.701 seconds, Fetched: 4 row(s)
hive> █
```

⋮ RAM 6.71 GB CPU 0.31% Disk: 7.71 GB used (limit 1006.85 GB)

## Query 2:

Finds the most-played content categories.

```
SELECT
    c.category,
    COUNT(*) AS play_count
FROM fact_user_activity f
JOIN dim_content c ON f.content_id = c.content_id
WHERE f.action = 'play'
AND f.year = 2025 AND f.month = 3 -- Example filter
GROUP BY c.category
ORDER BY play_count DESC
LIMIT 10;
```

```

MapReduce Jobs Launched:
Stage-Stage-2: Map: 1  Reduce: 1    Cumulative CPU: 4.12 sec  |
Stage-Stage-3: Map: 1  Reduce: 1    Cumulative CPU: 2.37 sec  |
Total MapReduce CPU Time Spent: 6 seconds 490 msec
OK
Jazz      28
Rock      8
Podcast   6
News      2
Time taken: 39.995 seconds, Fetched: 4 row(s)
hive> █

```

## HIVE DDL For Raw Tables and External Tables and Partitioning in HIVE

### +DDL For Star Schema + DDL For Transformation + Query 1 HIVE DDL + Query 2 DDL + Query 3 DDL

```

hive> CREATE EXTERNAL TABLE IF NOT EXISTS raw_user_logs (
  > user_id INT,
  > content_id INT,
  > action STRING,
  > timestamp STRING,
  > device STRING,
  > region STRING,
  > session_id STRING
  > )
  > PARTITIONED BY (year INT, month INT, day INT)
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY ','

```

> STORED AS TEXTFILE

> LOCATION '/raw/logs';

NoViableAltException(287@[[]])

at org.apache.hadoop.hive.ql.parse.HiveParser.columnNameTypeOrPKOrFK(HiveParser.java:33341)

at org.apache.hadoop.hive.ql.parse.HiveParser.columnNameTypeOrPKOrFKList(HiveParser.java:29513)

at org.apache.hadoop.hive.ql.parse.HiveParser.createTableStatement(HiveParser.java:6175)

at org.apache.hadoop.hive.ql.parse.HiveParser.ddlStatement(HiveParser.java:3808)

at org.apache.hadoop.hive.ql.parse.HiveParser.execStatement(HiveParser.java:2382)

at org.apache.hadoop.hive.ql.parse.HiveParser.statement(HiveParser.java:1333)

at org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:208)

at org.apache.hadoop.hive.ql.parse.ParseUtils.parse(ParseUtils.java:77)

at org.apache.hadoop.hive.ql.parse.ParseUtils.parse(ParseUtils.java:70)

at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:468)

at org.apache.hadoop.hive.ql.Driver.compileInternal(Driver.java:1317)

at org.apache.hadoop.hive.ql.Driver.runInternal(Driver.java:1457)

at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1237)

at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1227)

at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:233)

at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:184)

at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:403)

at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:821)

at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:759)

at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:686)

at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)

at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)

at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)

at java.lang.reflect.Method.invoke(Method.java:498)

at org.apache.hadoop.util.RunJar.run(RunJar.java:221)

at org.apache.hadoop.util.RunJar.main(RunJar.java:136)

FAILED: ParseException line 5:4 cannot recognize input near 'timestamp' 'STRING' ';' in column name or primary key or foreign key

hive> # hive



SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/opt/hive/lib/log4j-slf4j-impl-2.17.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/opt/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/opt/hive/lib/hive-common-2.3.10.jar!/hive-log4j2.properties Async: true

Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS raw_user_logs (
```

```
> user_id INT,
```

```
> content_id INT,
```

```
> action STRING,
```

```
> `timestamp` STRING, -- Use backticks to avoid conflicts
```

```
> device STRING,
```

```
> region STRING,
```

```
> session_id STRING
```

```
> )
```

```
> PARTITIONED BY (year INT, month INT, day INT)
```

```
> ROW FORMAT DELIMITED
```

```
> FIELDS TERMINATED BY ','
```

```
> STORED AS TEXTFILE
```

```
> LOCATION '/raw/logs';
```

OK

Time taken: 2.912 seconds

```
hive> ALTER TABLE raw_user_logs ADD PARTITION (year=2025, month=03, day=11) LOCATION '/raw/logs/2025/03/11/';
```

OK

Time taken: 0.2 seconds

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS raw_content_metadata (
```

```
> content_id INT,
```

```
> title STRING,
```

```
> category STRING,  
> length INT,  
> artist STRING  
> )  
> ROW FORMAT DELIMITED  
> FIELDS TERMINATED BY ','  
> STORED AS TEXTFILE  
> LOCATION '/raw/metadata';
```

OK

Time taken: 0.059 seconds

```
hive> SHOW TABLES;
```

OK

raw\_content\_metadata

raw\_user\_logs

Time taken: 0.072 seconds, Fetched: 2 row(s)

```
hive> DESCRIBE FORMATTED raw_user_logs;
```

OK

#	col_name	data_type	comment
---	----------	-----------	---------

user_id	int	
---------	-----	--

content_id	int	
------------	-----	--

action	string	
--------	--------	--

timestamp	string	
-----------	--------	--

device	string	
--------	--------	--

region	string	
--------	--------	--

session_id	string	
------------	--------	--

# Partition Information

#	col_name	data_type	comment
---	----------	-----------	---------

year	int	
------	-----	--

month int

day int

# Detailed Table Information

Database: default

Owner: root

CreateTime: Tue Mar 11 12:31:37 UTC 2025

LastAccessTime: UNKNOWN

Retention: 0

Location: hdfs://hive:9820/raw/logs

Table Type: EXTERNAL\_TABLE

Table Parameters:

EXTERNAL TRUE

numFiles 1

numPartitions 1

numRows 0

rawDataSize 0

totalSize 9764

transient\_lastDdlTime 1741696297

# Storage Information

SerDe Library: org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe

InputFormat: org.apache.hadoop.mapred.TextInputFormat

OutputFormat: org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat

Compressed: No

Num Buckets: -1

Bucket Columns: []

Sort Columns: []

Storage Desc Params:

field.delim ,

serialization.format ,

Time taken: 0.272 seconds, Fetched: 45 row(s)

hive> DESCRIBE FORMATTED raw\_content\_metadata;

OK

# col_name	data_type	comment
------------	-----------	---------

content_id	int	
------------	-----	--

title	string	
-------	--------	--

category	string	
----------	--------	--

length	int	
--------	-----	--

artist	string	
--------	--------	--

#### # Detailed Table Information

Database: default

Owner: root

CreateTime: Tue Mar 11 12:32:34 UTC 2025

LastAccessTime: UNKNOWN

Retention: 0

Location: hdfs://hive:9820/raw/metadata

Table Type: EXTERNAL\_TABLE

#### Table Parameters:

EXTERNAL	TRUE
----------	------

numFiles	1
----------	---

totalSize	431
-----------	-----

transient_lastDdlTime	1741696354
-----------------------	------------

#### # Storage Information

SerDe Library: org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe

InputFormat: org.apache.hadoop.mapred.TextInputFormat

OutputFormat: org.apache.hadoop.hive ql.io.HiveIgnoreKeyTextOutputFormat

Compressed: No

Num Buckets: -1

Bucket Columns: []

Sort Columns: []

Storage Desc Params:

field.delim ,

serialization.format ,

Time taken: 0.053 seconds, Fetched: 33 row(s)

```
hive> SELECT * FROM raw_user_logs WHERE year=2025 AND month=3 AND day=11 LIMIT 10;
```

OK

NULL	NULL	action	timestamp	device	region	session_id	2025	3	11
133	1004	play	2023-09-01 00:00:00	desktop	US	sess_3	2025	3	11
110	1008	forward	2023-09-01 00:00:00	tablet	APAC	sess_24	2025	3	11
125	1004	skip	2023-09-01 00:00:00	mobile	US	sess_23	2025	3	11
110	1003	forward	2023-09-01 00:00:00	mobile	APAC	sess_22	2025	3	11
157	1002	forward	2023-09-01 00:00:00	tablet	EU	sess_4	2025	3	11
122	1005	skip	2023-09-01 00:00:00	tablet	EU	sess_1	2025	3	11
119	1002	forward	2023-09-01 00:00:00	desktop	APAC	sess_41	2025	3	11
181	1006	play	2023-09-01 00:00:00	mobile	US	sess_28	2025	3	11
196	1010	pause	2023-09-01 00:00:00	desktop	APAC	sess_33	2025	3	11

Time taken: 1.361 seconds, Fetched: 10 row(s)

```
hive> SELECT * FROM raw_content_metadata LIMIT 10;
```

OK

NULL	title	category	NULL	artist
1000	Title 1000	Jazz	290	Artist 1
1001	Title 1001	Jazz	149	Artist 7
1002	Title 1002	Jazz	179	Artist 9
1003	Title 1003	Rock	283	Artist 1
1004	Title 1004	Jazz	149	Artist 10
1005	Title 1005	Podcast	192	Artist 10
1006	Title 1006	Rock	148	Artist 5
1007	Title 1007	Jazz	232	Artist 3
1008	Title 1008	News	181	Artist 8

Time taken: 0.108 seconds, Fetched: 10 row(s)

```
hive> CREATE TABLE dim_users STORED AS PARQUET AS
```

```
> SELECT DISTINCT user_id FROM raw_user_logs;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = root\_20250311133055\_c1a40568-a797-45e4-8662-d3ab582d7121

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0001, Tracking URL = http://hive:8088/proxy/application\_1741695021252\_0001/

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0001

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2025-03-11 13:31:04,273 Stage-1 map = 0%, reduce = 0%

2025-03-11 13:31:09,362 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.86 sec

2025-03-11 13:31:14,519 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.57 sec

MapReduce Total cumulative CPU time: 4 seconds 570 msec

Ended Job = job\_1741695021252\_0001

Moving data to directory hdfs://hive:9820/user/hive/warehouse/dim\_users

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.57 sec HDFS Read: 18015 HDFS Write: 641 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 570 msec

OK

Time taken: 20.587 seconds

```
hive> CREATE TABLE dim_content STORED AS PARQUET AS
```

```
> SELECT DISTINCT content_id, title, category, artist
```

```
> FROM raw_content_metadata
```

```
> WHERE content_id IS NOT NULL;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = root\_20250311133138\_1611034c-5f9c-4a84-b817-273763027230

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0002, Tracking URL = http://hive:8088/proxy/application\_1741695021252\_0002/

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0002

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2025-03-11 13:31:43,348 Stage-1 map = 0%, reduce = 0%

2025-03-11 13:31:47,466 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.5 sec

2025-03-11 13:31:52,589 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.44 sec

MapReduce Total cumulative CPU time: 4 seconds 440 msec

Ended Job = job\_1741695021252\_0002

Moving data to directory hdfs://hive:9820/user/hive/warehouse/dim\_content

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.44 sec HDFS Read: 9184 HDFS Write: 1023 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 440 msec

OK

Time taken: 15.544 seconds

## HIVE DDL For Star Schema:

```
hive> CREATE TABLE dim_users STORED AS PARQUET AS
```

```
> SELECT DISTINCT user_id FROM raw_user_logs;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = root\_20250311133055\_c1a40568-a797-45e4-8662-d3ab582d7121

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0001, Tracking URL = http://hive:8088/proxy/application\_1741695021252\_0001/

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0001

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2025-03-11 13:31:04,273 Stage-1 map = 0%, reduce = 0%

2025-03-11 13:31:09,362 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.86 sec

2025-03-11 13:31:14,519 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.57 sec

MapReduce Total cumulative CPU time: 4 seconds 570 msec

Ended Job = job\_1741695021252\_0001

Moving data to directory hdfs://hive:9820/user/hive/warehouse/dim\_users

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.57 sec HDFS Read: 18015 HDFS Write: 641 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 570 msec

OK

Time taken: 20.587 seconds

hive> CREATE TABLE dim\_content STORED AS PARQUET AS

```
> SELECT DISTINCT content_id, title, category, artist
```

```
> FROM raw_content_metadata
```

```
> WHERE content_id IS NOT NULL;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.



Query ID = root\_20250311133138\_1611034c-5f9c-4a84-b817-273763027230

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0002, Tracking URL = http://hive:8088/proxy/application\_1741695021252\_0002/

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0002

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2025-03-11 13:31:43,348 Stage-1 map = 0%, reduce = 0%

2025-03-11 13:31:47,466 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.5 sec

2025-03-11 13:31:52,589 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.44 sec

MapReduce Total cumulative CPU time: 4 seconds 440 msec

Ended Job = job\_1741695021252\_0002

Moving data to directory hdfs://hive:9820/user/hive/warehouse/dim\_content

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.44 sec HDFS Read: 9184 HDFS Write: 1023 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 440 msec

OK

Time taken: 15.544 seconds

hive> CREATE TABLE dim\_sessions STORED AS PARQUET AS

```
> SELECT DISTINCT session_id, device, region
```

```
> FROM raw_user_logs;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = root\_20250311133211\_bd71f6da-c762-467c-9b07-fb89e1b6f5c2

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0003, Tracking URL = http://hive:8088/proxy/application\_1741695021252\_0003/

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0003

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2025-03-11 13:32:15,178 Stage-1 map = 0%, reduce = 0%

2025-03-11 13:32:19,280 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.4 sec

2025-03-11 13:32:24,401 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.43 sec

MapReduce Total cumulative CPU time: 3 seconds 430 msec

Ended Job = job\_1741695021252\_0003

Moving data to directory hdfs://hive:9820/user/hive/warehouse/dim\_sessions

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.43 sec HDFS Read: 19039 HDFS Write: 1366 SUCCESS

Total MapReduce CPU Time Spent: 3 seconds 430 msec

OK

Time taken: 14.523 seconds

```
hive> CREATE TABLE fact_user_activity (
```

```
> user_id INT,
```

```
> content_id INT,
```

```
> session_id STRING,
```

```
> action STRING,
```

```
> event_timestamp STRING,
```

```
> year INT,
```

```
> month INT,
```

```
> day INT
```

```
> ) STORED AS PARQUET;
```

OK

Time taken: 0.05 seconds

```
hive> INSERT OVERWRITE TABLE fact_user_activity
```

```
> SELECT user_id, content_id, session_id, action, `timestamp`, year, month, day
```

```
> FROM raw_user_logs;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = root\_20250311133247\_87b21063-ce00-4bcf-9bf4-0c458bda021c

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1741695021252\_0004, Tracking URL = http://hive:8088/proxy/application\_1741695021252\_0004/

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0004

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2025-03-11 13:32:52,308 Stage-1 map = 0%, reduce = 0%

2025-03-11 13:32:56,406 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.92 sec

MapReduce Total cumulative CPU time: 1 seconds 920 msec

Ended Job = job\_1741695021252\_0004

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to directory hdfs://hive:9820/user/hive/warehouse/fact\_user\_activity/.hive-staging\_hive\_2025-03-11\_13-32-47\_265\_2347698501231064156-1/-ext-10000

Loading data to table default.fact\_user\_activity

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 1.92 sec HDFS Read: 15079 HDFS Write: 2912 SUCCESS

Total MapReduce CPU Time Spent: 1 seconds 920 msec

OK

Time taken: 10.483 seconds

```
hive> SELECT * FROM fact_user_activity LIMIT 10;
```

OK

SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".

SLF4J: Defaulting to no-operation (NOP) logger implementation

SLF4J: See <http://www.slf4j.org/codes.html#StaticLoggerBinder> for further details.

	NULL	NULL	session_id	action	timestamp	2025	3	11
133	1004	sess_3	play	2023-09-01 00:00:00	2025	3	11	
110	1008	sess_24	forward	2023-09-01 00:00:00	2025	3	11	
125	1004	sess_23	skip	2023-09-01 00:00:00	2025	3	11	
110	1003	sess_22	forward	2023-09-01 00:00:00	2025	3	11	
157	1002	sess_4	forward	2023-09-01 00:00:00	2025	3	11	
122	1005	sess_1	skip	2023-09-01 00:00:00	2025	3	11	
119	1002	sess_41	forward	2023-09-01 00:00:00	2025	3	11	
181	1006	sess_28	play	2023-09-01 00:00:00	2025	3	11	
196	1010	sess_33	pause	2023-09-01 00:00:00	2025	3	11	

Time taken: 0.105 seconds, Fetched: 10 row(s)

```
hive> SELECT * FROM dim_users LIMIT 10;
```

OK

NULL

100

103

104

105

106

107

108

109

110

Time taken: 0.069 seconds, Fetched: 10 row(s)

```
hive> SELECT * FROM dim_content LIMIT 10;
```

OK

1000 Title 1000 Jazz Artist 1

1001 Title 1001 Jazz Artist 7

1002 Title 1002 Jazz Artist 9

1003 Title 1003 Rock Artist 1

1004 Title 1004 Jazz Artist 10

1005 Title 1005 Podcast Artist 10

1006 Title 1006 Rock Artist 5

1007 Title 1007 Jazz Artist 3

1008 Title 1008 News Artist 8

1009 Title 1009 Jazz Artist 9

Time taken: 0.082 seconds, Fetched: 10 row(s)

```
hive> SELECT * FROM dim_sessions LIMIT 10;
```

OK

sess\_10 desktop APAC

sess\_11 desktop APAC

sess\_14 desktop APAC

sess\_16 desktop APAC

sess\_2 desktop APAC

sess\_21 desktop APAC

sess\_22 desktop APAC

sess\_3 desktop APAC

sess\_30 desktop APAC

sess\_33 desktop APAC

Time taken: 0.072 seconds, Fetched: 10 row(s)

#### **Partitioned table:**

```
CREATE EXTERNAL TABLE IF NOT EXISTS user_logs_partitioned (
```

```
  user_id STRING,
```

```
  event_type STRING,
```

```
  event_timestamp STRING,
```

```
  event_details STRING
```

```
)
```

```
PARTITIONED BY (year INT, month INT, day INT)
```

```
ROW FORMAT DELIMITED
```

FIELDS TERMINATED BY ','

STORED AS TEXTFILE

LOCATION '/raw/logs/';

### **HIVE DDL For Transformation:**

```
hive> INSERT OVERWRITE TABLE fact_user_activity
```

```
> SELECT user_id, content_id, session_id, action, `timestamp`, year, month, day
```

```
> FROM raw_user_logs;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = root\_20250311133247\_87b21063-ce00-4bcf-9bf4-0c458bda021c

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1741695021252\_0004, Tracking URL = http://hive:8088/proxy/application\_1741695021252\_0004/

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0004

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2025-03-11 13:32:52,308 Stage-1 map = 0%, reduce = 0%

2025-03-11 13:32:56,406 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.92 sec

MapReduce Total cumulative CPU time: 1 seconds 920 msec

Ended Job = job\_1741695021252\_0004

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to directory hdfs://hive:9820/user/hive/warehouse/fact\_user\_activity/.hive-staging\_hive\_2025-03-11\_13-32-47\_265\_2347698501231064156-1/-ext-10000

Loading data to table default.fact\_user\_activity

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 1.92 sec HDFS Read: 15079 HDFS Write: 2912 SUCCESS

Total MapReduce CPU Time Spent: 1 seconds 920 msec

OK

Time taken: 10.483 seconds

## Query 1 DDL:

```
hive> SELECT
```

```
> f.year,
```

```
> f.month,
```

```
> s.region,
```

```
> COUNT(DISTINCT f.user_id) AS monthly_active_users
```

```
> FROM fact_user_activity f
```

```
> JOIN dim_sessions s ON f.session_id = s.session_id
```

```
> WHERE f.year = 2025 -- Example filter
```

```
> GROUP BY f.year, f.month, s.region
```

```
> ORDER BY f.year, f.month, monthly_active_users DESC;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = root\_20250311142406\_3fa5adc6-3b22-4304-a693-73eea3ebe2b7

Total jobs = 2

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/opt/hive/lib/log4j-slf4j-impl-2.17.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/opt/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

2025-03-11 14:24:16 Starting to launch local task to process map join; maximum memory = 477626368

SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".

SLF4J: Defaulting to no-operation (NOP) logger implementation

SLF4J: See <http://www.slf4j.org/codes.html#StaticLoggerBinder> for further details.

2025-03-11 14:24:17 Dump the side-table for tag: 1 with group count: 50 into file: file:/tmp/root/f280238c-aeba-4b03-9fa3-42d8adb61222/hive\_2025-03-11\_14-24-06\_316\_3053109984737745539-1/-local-10006/HashTable-Stage-2/MapJoin-mapfile01--.hashtable

2025-03-11 14:24:17 Uploaded 1 File to: file:/tmp/root/f280238c-aeba-4b03-9fa3-42d8adb61222/hive\_2025-03-11\_14-24-06\_316\_3053109984737745539-1/-local-10006/HashTable-Stage-2/MapJoin-mapfile01--.hashtable (2585 bytes)

2025-03-11 14:24:17 End of local task; Time Taken: 1.486 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 2

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0005, Tracking URL =  
[http://hive:8088/proxy/application\\_1741695021252\\_0005/](http://hive:8088/proxy/application_1741695021252_0005/)

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0005

Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1

2025-03-11 14:24:24,837 Stage-2 map = 0%, reduce = 0%

2025-03-11 14:24:29,902 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.95 sec

2025-03-11 14:24:35,082 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 4.56 sec

MapReduce Total cumulative CPU time: 4 seconds 560 msec

Ended Job = job\_1741695021252\_0005

Launching Job 2 out of 2

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0006, Tracking URL =  
[http://hive:8088/proxy/application\\_1741695021252\\_0006/](http://hive:8088/proxy/application_1741695021252_0006/)



Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0006

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 1

2025-03-11 14:24:45,716 Stage-3 map = 0%, reduce = 0%

2025-03-11 14:24:48,797 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 0.91 sec

2025-03-11 14:24:52,896 Stage-3 map = 100%, reduce = 100%, Cumulative CPU 2.19 sec

MapReduce Total cumulative CPU time: 2 seconds 190 msec

Ended Job = job\_1741695021252\_0006

MapReduce Jobs Launched:

Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 4.56 sec HDFS Read: 16306 HDFS Write: 190 SUCCESS

Stage-Stage-3: Map: 1 Reduce: 1 Cumulative CPU: 2.19 sec HDFS Read: 6216 HDFS Write: 192 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 750 msec

OK

2025 3 US 80

2025 3 APAC 79

2025 3 EU 72

2025 3 region 0

Time taken: 48.701 seconds, Fetched: 4 row(s)

## Query 2 DDL

hive> SELECT

> c.category,

> COUNT(\*) AS play\_count

> FROM fact\_user\_activity f

> JOIN dim\_content c ON f.content\_id = c.content\_id

> WHERE f.action = 'play'

> AND f.year = 2025 AND f.month = 3 -- Example filter

> GROUP BY c.category

> ORDER BY play\_count DESC

> LIMIT 10;

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = root\_20250311142721\_971cebab-7e1c-406d-86a7-f1598c563ead

Total jobs = 2

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/opt/hive/lib/log4j-slf4j-impl-2.17.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/opt/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

2025-03-11 14:27:26 Starting to launch local task to process map join; maximum memory = 477626368

SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".

SLF4J: Defaulting to no-operation (NOP) logger implementation

SLF4J: See <http://www.slf4j.org/codes.html#StaticLoggerBinder> for further details.

2025-03-11 14:27:27 Dump the side-table for tag: 1 with group count: 11 into file: file:/tmp/root/f280238c-aeba-4b03-9fa3-42d8adb61222/hive\_2025-03-11\_14-27-21\_806\_1616347904249172090-1/-local-10006/HashTable-Stage-2/MapJoin-mapfile11--.hashtable

2025-03-11 14:27:27 Uploaded 1 File to: file:/tmp/root/f280238c-aeba-4b03-9fa3-42d8adb61222/hive\_2025-03-11\_14-27-21\_806\_1616347904249172090-1/-local-10006/HashTable-Stage-2/MapJoin-mapfile11--.hashtable (552 bytes)

2025-03-11 14:27:27 End of local task; Time Taken: 1.043 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 2

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0007, Tracking URL =  
[http://hive:8088/proxy/application\\_1741695021252\\_0007/](http://hive:8088/proxy/application_1741695021252_0007/)

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0007

Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1

2025-03-11 14:27:33,884 Stage-2 map = 0%, reduce = 0%

2025-03-11 14:27:37,990 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.7 sec

2025-03-11 14:27:42,089 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 4.12 sec

MapReduce Total cumulative CPU time: 4 seconds 120 msec

Ended Job = job\_1741695021252\_0007

Launching Job 2 out of 2

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1741695021252\_0008, Tracking URL =  
[http://hive:8088/proxy/application\\_1741695021252\\_0008/](http://hive:8088/proxy/application_1741695021252_0008/)

Kill Command = /opt/hadoop/bin/hadoop job -kill job\_1741695021252\_0008

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 1

2025-03-11 14:27:52,621 Stage-3 map = 0%, reduce = 0%

2025-03-11 14:27:56,621 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 0.99 sec

2025-03-11 14:28:00,737 Stage-3 map = 100%, reduce = 100%, Cumulative CPU 2.37 sec

MapReduce Total cumulative CPU time: 2 seconds 370 msec

Ended Job = job\_1741695021252\_0008

MapReduce Jobs Launched:

Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 4.12 sec HDFS Read: 13890 HDFS Write: 191 SUCCESS

Stage-Stage-3: Map: 1 Reduce: 1 Cumulative CPU: 2.37 sec HDFS Read: 5751 HDFS Write: 167 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 490 msec

OK

Jazz 28

Rock 8

Podcast 6

News 2

Time taken: 39.995 seconds, Fetched: 4 row(s)

hive>

Query 3 DDL: